



Shri Vile Parle Kelavani Mandal's

DWARKADAS J. SANGHVI COLLEGE OF ENGINEERING

(Autonomous College Affiliated to the University of Mumbai)

NAAC Accredited with "A" Grade (CGPA : 3.18)



Experiment no: 1

Case Study

Name: Kresha Shah

SAP ID: 60009220080

Domain: Gaming

Topic: Flappy bird

Summary of the topic


The study explores the application of deep reinforcement learning in the context of the popular game Flappy Bird. It aims to understand how an agent can learn to play the game effectively using Q-learning and a convolutional neural network. The agent receives pixel information and the game score as input, attempting to navigate the challenges of the game environment. The study sheds light on the agent's ability to navigate the game environment, showcasing the potential of reinforcement learning in mastering complex gaming tasks like Flappy Bird. The research delves into the challenges associated with high-dimensional input and sparse rewards in the context of Flappy Bird, providing insights into the capabilities of reinforcement learning in a gaming scenario.

Agent

The agent used in this project for playing Flappy Bird is based on deep reinforcement learning, specifically utilizing a Deep Q-Network (DQN). Here's a breakdown:

1. **Agent Type:** The agent is an artificial intelligence system trained using reinforcement learning techniques. It interacts with the Flappy Bird game environment and learns to make decisions (flapping or not flapping) in order to maximize its cumulative reward, which is tied to the game score.
2. **Architecture:** The architecture of the agent's neural network is a Convolutional Neural Network (CNN) known as the Deep Q-Network (DQN). The DQN takes pixel information from the game screens as input and produces Q-values for different actions as output.
 - The first layer is a convolutional layer with 32 filters of size 8x8, using a stride of 4.
 - The second layer is another convolutional layer with 64 filters of size 4x4, using a stride of 2.
 - The third layer is a convolutional layer with 64 filters of size 3x3 and a stride of 1.
 - Following the convolutional layers, there is a fully connected layer with 512 outputs.
 - The final layer is the output layer with a single output for each possible action.
3. **Training Approach:** The training utilizes a variant of Q-learning, where the Q-function is approximated by the DQN. The agent learns to associate actions with states by iteratively updating its Q-values based on the Bellman equation.
4. **Training Techniques:**
 - **Experience Replay:** Experiences (state, action, reward, next state) are stored in a replay memory. During training, batches of experiences are sampled uniformly from this memory, which helps in decorrelating experiences and improving the stability of training.
 - **ϵ -Greedy Exploration:** During training, the agent uses an ϵ -greedy approach to balance exploration and exploitation. It chooses a random action with probability ϵ and the action with the highest Q-value otherwise.
 - **Target Network:** To enhance stability during training, a target network (Q^*) is used, and it is updated periodically.
5. **Training Results:**
 - The trained DQN shows super-human performance, surpassing both a baseline agent (flapping every z frames) and human players in terms of the average and highest scores.
 - The agent's performance is evaluated on different difficulty levels (easy, medium, hard), showcasing its ability to generalize across various challenges in the game.
6. **Additional Experiment:**
 - The study includes an experiment where a pre-trained network on an easier difficulty level is fine-tuned on a harder difficulty. Results show that the DQN performs better

on both the trained and easier difficulties, demonstrating some transfer learning capabilities.



State

The state refers to the current representation of the game environment, encompassing the essential information required for the agent to make decisions. The state includes a sequence of frames from the Flappy Bird game, with each frame capturing the pixel input or screen capture at a specific time. Additionally, the state incorporates recent actions taken by the player, ensuring the agent has temporal information to understand the dynamics of the game. The state, denoted as s_t , is constructed as a tuple $(x_{t-\text{histLen}+1}, \alpha_{t-\text{histLen}+1}, \dots, x_{t-1}, a_{t-1}, x_t)$ where x_t represents the pixel input at time t , and histLen is a hyperparameter determining the number of recent frames to consider. This formulation enables the agent to navigate the challenges of the Flappy Bird environment by incorporating both spatial and temporal aspects of the game.

Reward

In the Flappy Bird reinforcement learning project, the reward system plays a crucial role in shaping the agent's learning process. The goal is to teach the agent to successfully navigate the game environment, and rewards serve as feedback mechanisms that guide the agent toward desirable behavior. The reward scheme includes three key components: `rewardAlive`, `rewardPipe`, and `rewardDead`.

1. **rewardAlive:** This reward is assigned for every frame the agent stays alive. It provides a continuous incentive for the agent to prolong its survival in the game. Given the sparse nature of other rewards, `rewardAlive` contributes to faster learning by encouraging actions that lead to sustained gameplay.
2. **rewardPipe:** Awarded when the agent successfully passes through a pipe, `rewardPipe` is a positive reinforcement tied to the primary objective of the game. It encourages the agent to learn the optimal timing for navigating through obstacles, ultimately contributing to achieving a higher game score.
3. **rewardDead:** This penalty is incurred when the agent's performance results in its demise. It serves as a negative reinforcement, discouraging actions that lead to failure. By associating a cost with unfavorable outcomes, `rewardDead` guides the agent to avoid actions that lead to premature termination of the game.

The combination of these rewards provides a balanced framework for the agent to learn effective strategies.

Value-> state, action

Values refer to the numerical estimates associated with different states and actions. These values play a crucial role in guiding the learning process of the agent. The value associated with a state represents the expected cumulative reward the agent anticipates when in that particular state. Similarly, the value associated with a specific action in a given state represents the expected cumulative reward of taking that action in that particular situation. The agent learns and updates these values through its interactions with the game environment, aiming to improve its decision-making abilities and overall performance in Flappy Bird.

State Values ($V(s)$): State values in Flappy Bird represent the expected cumulative reward an agent anticipates when in a particular game state. For example, a high state value might indicate that being at a certain height relative to the pipes and obstacles is advantageous. The agent learns these state values through its interactions with the environment, aiming to understand the desirability of different situations within the game. State values help the agent navigate the complex spatial dynamics of Flappy Bird and make decisions that lead to favorable outcomes.

Action Values ($Q(s, a)$): Action values, or Q-values, represent the expected cumulative reward of taking a specific action (a) in a given state (s) and then following a certain policy. In Flappy Bird, actions include decisions like flapping or not flapping. Q-values enable the agent to assess the desirability of each action in different game situations. Through reinforcement learning algorithms like Q-learning, the agent refines its understanding of which actions lead to higher rewards in specific states, improving its overall performance.

Policy -> deterministic, stochastic

Flappy Bird employs a deterministic policy, where the agent's actions are based on predefined rules, with binary choices of either flapping or doing nothing in response to the game environment.

1. **Deterministic Policy:** The deterministic aspect of the policy implies that the agent's actions are not purely random but are influenced by a set of rules or a function. In this case, the agent has two deterministic actions it can take: either it flaps ($a = 1$) or does nothing ($a = 0$). This deterministic component provides a structured approach to decision-making, allowing the agent to make intentional choices based on its learning from the training data.
2. **Stochastic Policy:** In addition to deterministic actions, stochasticity is introduced to the policy. Stochastic elements involve randomness or probability in decision-making. For example, during training, the agent may not always choose the same action in a given state. Instead, it introduces an element of randomness in selecting its actions. This stochasticity can be essential for exploration, especially in the early stages of training when the agent is learning and needs to explore the range of possible actions to discover optimal strategies.

Model -> based, free

1. **Model-Based Learning:**

Definition: In model-based learning, the agent constructs an explicit model of the environment, which includes an understanding of the transition dynamics between states, the probabilities of transitioning between states, and the associated rewards.

2. **Model-Free Learning:**

Definition: Model-free learning, on the other hand, does not involve creating an explicit model of the environment. Instead, the agent learns directly from interacting with the environment and observing the consequences of its actions.

Now, regarding Flappy Bird:

- Flappy Bird employs a **model-free learning approach**. The agent learns to play the game by directly interacting with the environment, making decisions based on pixel information and the score, without explicitly modeling the transition dynamics or probabilities.
- Flappy Bird uses a variant of the Q-learning algorithm, a classic model-free reinforcement learning technique. The Q-function is approximated using a deep neural network, indicating a model-free, value-based learning approach.

Exploration

The terms "exploration" and "exploitation" refer to the agent's strategy in selecting actions.

- **Exploration:** This involves the agent trying new actions to discover more about the environment and improve its understanding of the optimal strategy. Exploratory actions may not necessarily be the best in terms of immediate rewards, but they contribute to the agent's learning process.

Exploitation

- **Exploitation:** This involves the agent selecting actions that it believes are currently the best based on its current knowledge or learned policy. Exploitative actions aim to maximize short-term rewards based on the agent's existing understanding of the environment.

In the case of Flappy Bird and similar reinforcement learning scenarios:

- During training, the agent (Flappy Bird playing algorithm) often incorporates an exploration-exploitation strategy. This is typically achieved using an epsilon-greedy approach, where with probability ϵ (epsilon), the agent selects a random action (exploration), and with probability $1-\epsilon$, it selects the action it believes to be the best based on its current learned policy (exploitation).
- Initially, the agent may heavily explore the action space to gain a diverse set of experiences and learn about the consequences of different actions. As training progresses, the balance may shift more toward exploitation as the agent refines its policy based on accumulated knowledge.

So, in the context of Flappy Bird training, the agent does engage in both exploration and exploitation to iteratively improve its understanding of the game dynamics and optimize its decision-making policy.

References:

https://cs229.stanford.edu/proj2015/362_report.pdf

<https://www.toptal.com/deep-learning/pytorch-reinforcement-learning-tutorial>

https://www.researchgate.net/publication/335969576_Reinforcement_Learning_and_Neuroevolution_in_Flappy_Bird_Game

Topic 2: Interactive story telling

Summary of the topic:

Interactive storytelling using reinforcement learning (RL) combines the narrative depth of traditional storytelling with the dynamic responsiveness of artificial intelligence. By employing RL algorithms, interactive storytelling systems can adapt narratives based on user actions and preferences, creating personalized and engaging experiences. These systems often model the story world as a Markov Decision Process (MDP), where the actions of characters and events unfold based on the current state and previous interactions. Through exploration and exploitation, RL agents learn optimal decision-making policies, allowing them to craft narratives that resonate with individual users while maintaining coherence and narrative flow. As users interact with the story, the RL agent continuously learns and adapts, shaping the narrative trajectory in real-time to provide immersive and captivating storytelling experiences tailored to each user's preferences and choices.

Agent

The agent used in the paper "Interactive Storytelling with Deep Reinforcement Learning" employs deep reinforcement learning techniques to generate interactive narratives. Here's an overview:

1. **Agent Type:** The agent is an artificial intelligence system trained with reinforcement learning methods. It interacts with a simulated environment representing a storytelling domain and learns to make decisions to progress the story, aiming to maximize predefined rewards associated with engaging and coherent narratives.

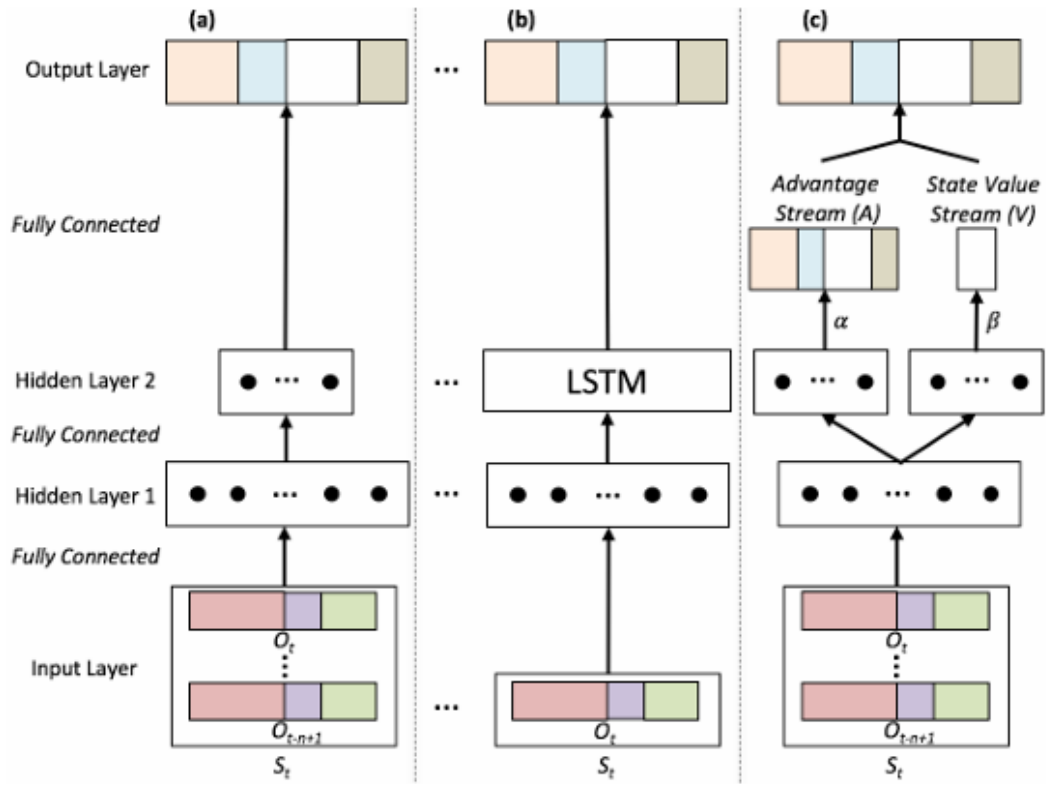
2. **Architecture:** The agent's neural network architecture consists of a deep reinforcement learning model, possibly utilizing variants of deep Q-learning or policy gradient methods. The specifics of the neural network architecture, such as the number of layers and types of layers, are not explicitly mentioned in the provided information.

3. **Training Approach:** The agent is trained using reinforcement learning algorithms tailored for sequential decision-making tasks, where it learns to generate story elements based on interactions with the environment and feedback received. The training process involves updating the parameters of the neural network to improve the agent's storytelling capabilities over time.

4. Training Techniques:

- Experience Replay: The agent utilizes experience replay to store and sample past interactions with the environment, facilitating more efficient learning by breaking correlations between sequential experiences.
- Exploration and Exploitation: The agent balances exploration and exploitation during training to discover new narrative possibilities while also exploiting learned strategies to maximize rewards.
- Reward Design: Rewards are designed to incentivize the generation of engaging and coherent narratives, providing feedback to the agent on the quality of its storytelling decisions.

5. **Training Results:** The paper likely presents results demonstrating the effectiveness of the trained agent in generating interactive stories compared to baseline methods or human-authored narratives. Metrics such as user engagement, coherence, and entertainment value may be used to evaluate the quality of the generated narratives.



Reward

1. **Reward Mechanism:** The reward system is designed to encourage the agent to produce narratives that are engaging and coherent. It provides feedback to the agent based on the quality of the generated narratives, with the goal of maximizing cumulative rewards over the course of storytelling.
2. **Calculation of Reward:** The paper outlines the criteria used to calculate the reward for each generated narrative. While specific details may vary depending on the implementation, the reward calculation typically considers factors such as:
 - **Narrative Coherence:** The extent to which the events and characters in the story follow a logical sequence and adhere to established storytelling conventions. Higher coherence leads to higher rewards.
 - **User Engagement:** The level of interest and involvement elicited from users interacting with the narrative. Engaging narratives that captivate users result in higher rewards.
 - **Story Quality:** The overall quality of the narrative in terms of its originality, creativity, and emotional impact. Well-crafted stories with compelling themes and characters receive higher rewards.
 - **Alignment with Objectives:** The degree to which the narrative fulfills specific objectives or constraints defined by the storytelling task. Aligning with these objectives leads to higher rewards.
 - **User Feedback:** Direct feedback from users or evaluators regarding their satisfaction with the narrative can also influence the reward calculation.
3. **Training Process:** During training, the agent generates multiple narratives and receives rewards based on the evaluation of each narrative against the predefined criteria. These rewards are used to update the parameters of the agent's neural network through reinforcement learning algorithms, such as Q-learning or policy gradients.
4. **Optimization Objective:** The agent's objective is to learn a policy that maximizes the expected cumulative reward over time. By iteratively adjusting its behavior based on received rewards, the agent learns to generate narratives that are more engaging and coherent, leading to improved performance.

Value

1. State Value:

- In the context of interactive storytelling, a state represents a particular configuration or context within the narrative generation process. This could include the current plot progression, character interactions, setting details, and other story-related elements.
- The state value, denoted as $V(s)$, quantifies the expected cumulative reward that the agent can achieve starting from a particular state s and following its policy to generate the rest of the narrative.
- The state value serves as a measure of the desirability of a given state in terms of its potential to lead to favorable outcomes, such as engaging narratives or user satisfaction.

2. Action Value:

- Actions in interactive storytelling refer to the decisions made by the agent at each step of the narrative generation process. These decisions could involve introducing new story elements, developing character arcs, resolving conflicts, or incorporating user preferences.
- The action value, denoted as $Q(s, a)$, represents the expected cumulative reward that the agent can achieve by taking a specific action a in a given state s and following its policy thereafter.
- The action value provides insight into the potential outcomes associated with different narrative choices, allowing the agent to prioritize actions that are likely to lead to higher rewards or user engagement.

Policy

1. Policy:

- The policy represents the agent's approach to selecting actions (storytelling decisions) based on the current state of the narrative and any available contextual information.
- In interactive storytelling, the policy dictates how the agent responds to user inputs, narrative constraints, and other dynamic factors to produce coherent and engaging stories.
- The goal of the policy is to maximize the expected cumulative reward or user satisfaction over the course of the narrative generation process, guiding the agent towards desirable storytelling outcomes.

2. Types of Policy: a. Deterministic Policy:

- A deterministic policy specifies a one-to-one mapping between states and actions, meaning that for any given state, the policy always selects the same action.
- In the context of interactive storytelling, a deterministic policy might involve predefined rules or algorithms that dictate how the agent responds to specific narrative situations or user interactions.
- While deterministic policies can be straightforward to implement and interpret, they may lack flexibility in handling uncertain or complex narrative scenarios.

b. Stochastic Policy:

- A stochastic policy selects actions probabilistically, meaning that for a given state, the policy may choose different actions with certain probabilities.
- Stochastic policies introduce randomness into the decision-making process, allowing the agent to explore different narrative possibilities and adapt to changing environments or user preferences.
- In interactive storytelling, a stochastic policy might incorporate randomness to inject variability and creativity into the narrative generation process, leading to more diverse and engaging storylines.
- Stochastic policies can facilitate exploration of the narrative space, potentially uncovering novel storylines or interactive elements that would not be discovered with a deterministic approach.

In the context of the paper, the specific type of policy employed by the agent for interactive storytelling is not explicitly mentioned. However, given the creative and dynamic nature of narrative generation, it is likely that a combination of deterministic and stochastic elements is utilized to balance structure and flexibility in storytelling.

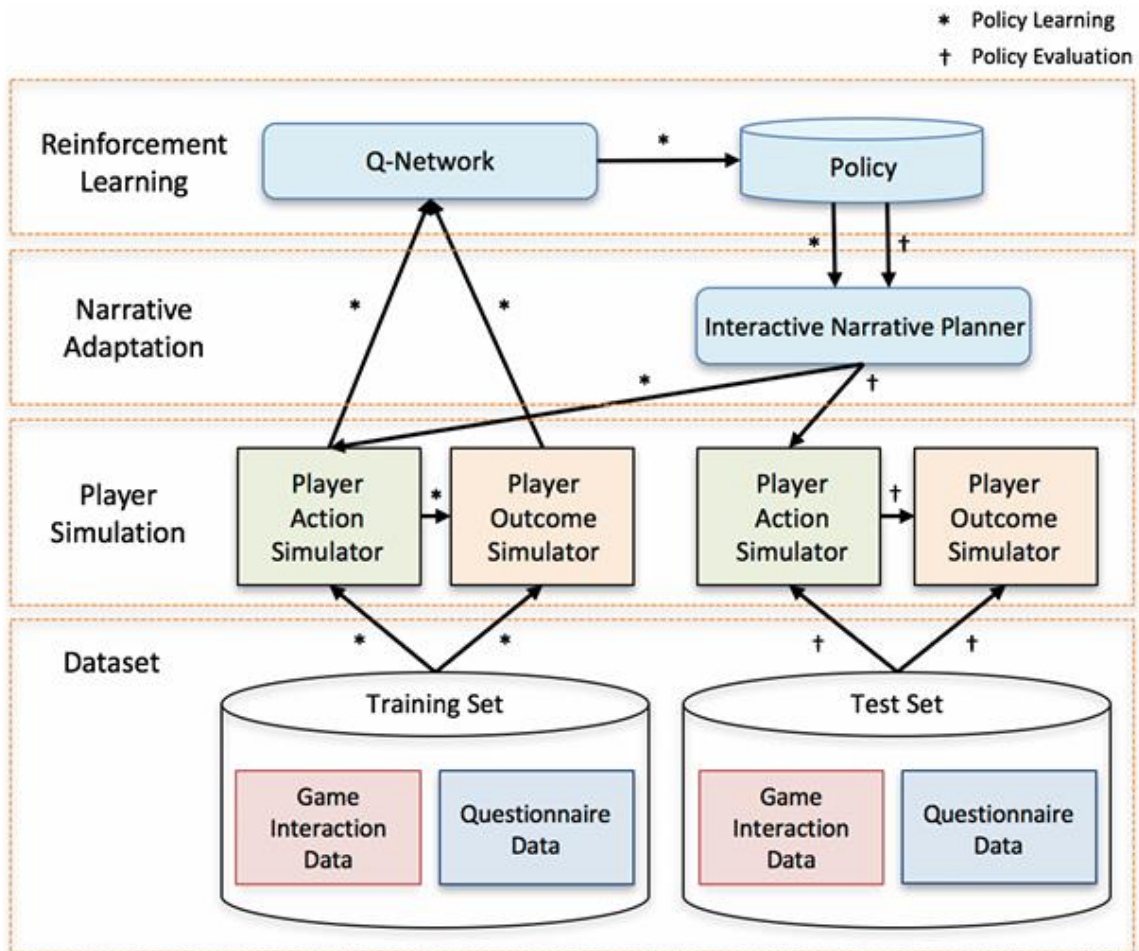


Figure 2. Q-network based deep RL interactive narrative personalization framework.

Model

The approach to narrative generation using deep reinforcement learning can be characterized as model-based. Here's how the concepts of model-based and model-free learning apply to the context of the paper:

1. Model-Based Learning:

- Model-based learning involves constructing an explicit model of the environment, including the dynamics of state transitions, rewards, and possible actions.
- In the context of narrative generation, a model-based approach would entail building a structured representation of the narrative space, including story elements, character interactions, and possible story trajectories.
- The agent would use this internal model to simulate and predict the consequences of different storytelling decisions, enabling it to plan and optimize narrative sequences based on expected outcomes.
- Model-based learning approaches often require accurate modeling of the environment dynamics, which may involve domain-specific knowledge and assumptions about narrative structure and coherence.

2. Model-Free Learning:

- Model-free learning, on the other hand, does not involve explicitly constructing a model of the environment. Instead, the agent learns directly from interactions with the environment, without explicitly modeling state transitions or rewards.
- In narrative generation, a model-free approach would focus on learning optimal storytelling policies through trial and error, without explicitly representing the underlying narrative structure or dynamics.
- Model-free methods, such as Q-learning or policy gradient methods, rely on estimating value functions or policies directly from observed experiences, without relying on an explicit model of the environment.
- Model-free learning approaches are often more flexible and adaptable, as they do not require detailed knowledge of the environment dynamics. However, they may require more extensive exploration and training to discover effective storytelling strategies.

In the context of the paper, the approach to interactive storytelling using deep reinforcement learning is more aligned with model-based learning. The agent learns to generate narratives by constructing an internal model of the narrative space and using it to simulate and optimize storytelling decisions.

Exploration vs Exploitation

The approach used involves a combination of exploration and exploitation strategies to generate engaging narratives. Here's how exploration and exploitation are implemented in the context of the paper:

1. Exploration:

- Exploration refers to the process of trying out different actions or narrative choices to gather information about the environment and discover potentially beneficial strategies.
- In the context of narrative generation, exploration involves generating diverse storylines, character interactions, and plot developments to explore the range of possibilities within the narrative space.
- The agent employs exploration strategies to sample novel story trajectories, introducing variability and creativity into the generated narratives.
- Exploration is essential for discovering new narrative elements, character behaviors, and story outcomes, which may lead to more engaging and unexpected storytelling experiences.

2. Exploitation:

- Exploitation involves leveraging known strategies or narrative elements to maximize the quality of the generated narratives based on existing knowledge and experiences.
- In narrative generation, exploitation entails selecting narrative choices or actions that are expected to lead to desirable story outcomes or align with user preferences.
- The agent exploits learned policies, preferences, and storytelling conventions to generate coherent and compelling storylines, drawing from past experiences and successful narrative sequences.
- Exploitation is crucial for refining and optimizing the generated narratives, focusing on high-quality storytelling elements and maximizing user satisfaction.

References:

[wang-ijcai-2017.pdf \(ncsu.edu\)](#)

[\(PDF\) Interactive Learning via Digital Storytelling in Teaching and Learning \(researchgate.net\)](#)

[Interactive storytelling for children: A case-study of design and development considerations for ethical conversational AI - ScienceDirect](#)

Topic 3: Language Learner

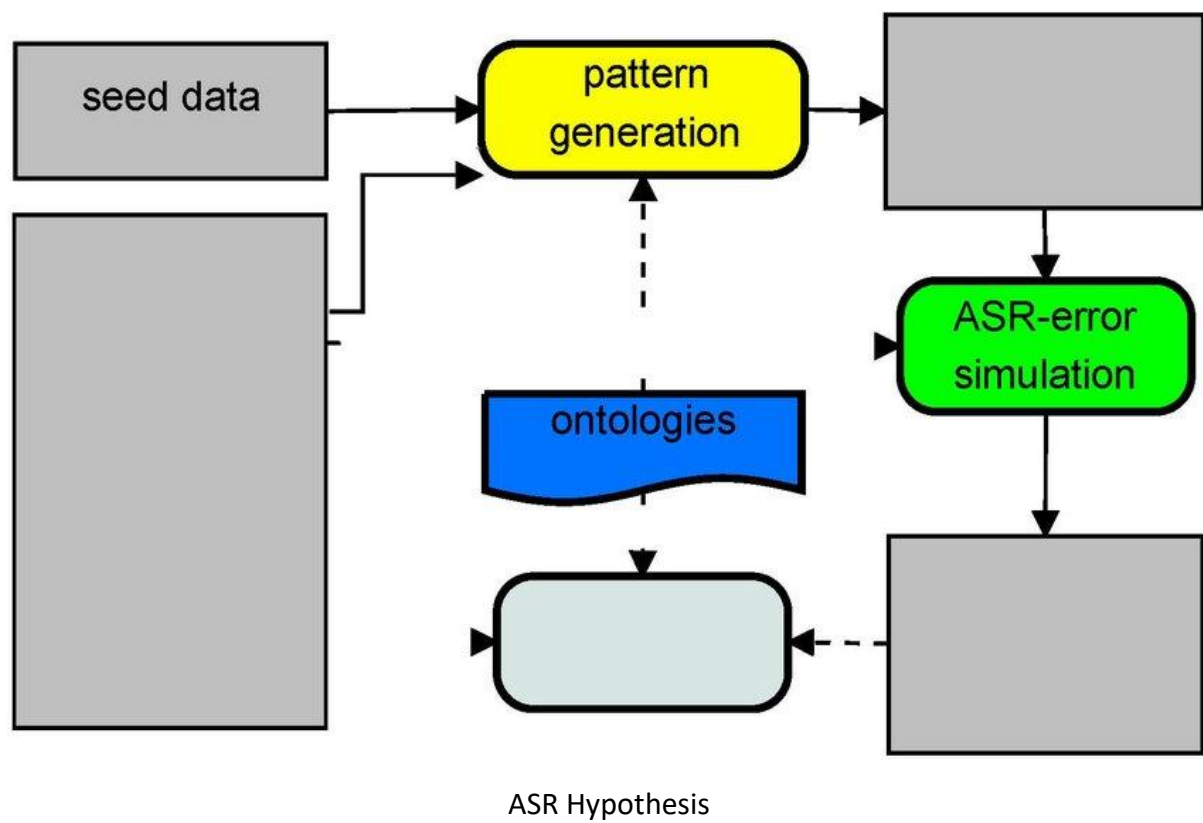
Summary:

Reinforcement learning (RL) can significantly enhance Robust Spoken Language Understanding (SLU) by facilitating RL-based Value Error Recovery mechanisms. RL enables SLU systems to learn from errors encountered during interaction with users, allowing them to adapt and improve over time. By employing RL, SLU systems can identify errors in speech recognition or understanding, assess the impact of potential actions (such as error correction strategies), and select the most effective course of action to recover from errors. This iterative learning process enables SLU systems to become more robust, resilient, and capable of handling diverse and challenging real-world scenarios. RL-based Value Error Recovery mechanisms offer a dynamic and adaptive approach to error handling in SLU systems, leading to improved performance, user satisfaction, and overall system reliability.

Agent

The agent used in this proposed method is an RL-based training algorithm designed to guide the training of a slot tagging model on Automatic Speech Recognition (ASR) hypotheses. Here's how it works:

1. **Agent Type:** The agent is an RL-based algorithm that interacts with the environment, which in this case is the slot tagging model trained on ASR hypotheses.
2. **Working:**
 - During the training of the slot tagging model, the RL-based algorithm provides guidance by considering the output of the slot tagging model and comparing it with the manually transcribed labels.
 - It samples tag sequences from the slot tagging model using beam search, producing multiple sets of act(slot=value) triplets.
 - These triplets are then corrected by a value error recovery (VER) module, which refines wrong values based on a pre-defined domain ontology.
 - The RL-based algorithm evaluates the correctness of the predicted triplets at both the triplet-level and utterance-level, calculating a reward based on the false-positive and false-negative predictions.
 - The model is then optimized using policy gradient descent to maximize the expected cumulative rewards, updating the parameters of the slot tagging model.
3. **Training Approach:** The RL-based algorithm utilizes policy gradient descent to optimize the slot tagging model, considering the rewards obtained from evaluating the correctness of the predicted triplets.
4. **Training Procedure:**
 - Pre-training: The slot tagging model is initially trained using manually transcribed labels to minimize a negative log-likelihood loss.
 - RL-training: ASR hypotheses, coupled with unaligned labels, are then used for adaptive training. The RL-based algorithm samples tag sequences, corrects them using the VER module, and optimizes the model parameters using policy gradient descent to maximize the expected cumulative rewards.



State

The "state" refers to the current representation of the environment or system being modeled. In this case, the state can be described as follows:

1. **ASR Hypotheses:** The input to the system consists of Automatic Speech Recognition (ASR) hypotheses, represented as sequences of words or tokens generated by an ASR system from spoken utterances. These ASR hypotheses serve as the primary source of information for the spoken language understanding task.
2. **Semantic Representation:** Along with the ASR hypotheses, the system also considers semantic representations, typically in the form of act(slot=value) triplets. These triplets annotate the ASR hypotheses with semantic information, indicating the intended meaning or intent behind the spoken utterances.
3. **Slot Tagging Model Output:** The output of the slot tagging model, which predicts slot tags (e.g., inform, deny, request) for each word or token in the ASR hypotheses. These predicted slot tags provide additional context about the semantic content of the spoken utterances.
4. **Value Error Recovery (VER) Module Output:** The corrected semantic triplets generated by the Value Error Recovery (VER) module. This module refines the semantic representations obtained from the ASR hypotheses by correcting errors in the predicted slot-value pairs.
5. **Reward Signals:** The rewards obtained during the evaluation of the predicted semantic triplets, which indicate the correctness of the predictions at both the triplet-level and utterance-level. These rewards influence the training of the slot tagging model through reinforcement learning.

Reward

In the RL-based training framework for robust spoken language understanding described in the paper, the reward function plays a crucial role in guiding the learning process. The reward is calculated based on the correctness of the predicted semantic triplets at both the triplet-level and utterance-level. Here's how the reward is considered and calculated:

1. **Triplet-level Reward**: At the triplet level, the reward function evaluates the correctness of the predicted semantic triplets (act(slot=value) triplets) generated by the system compared to the ground truth annotations. Specifically, the reward function penalizes false positives (FP) and false negatives (FN) of the predicted triplets. False positives occur when the system incorrectly predicts a slot-value pair that is not present in the ground truth annotations, while false negatives occur when the system fails to predict a slot-value pair that is present in the ground truth annotations. The reward is calculated as follows:

$$R_{\text{triplet}} = 1 - \text{FP}(y, \tilde{y}_k) + \text{FN}(y, \tilde{y}_k)$$

Where:

- $\text{FP}(y, \tilde{y}_k)$ is the number of false positives between the predicted triplet (\tilde{y}_k) and the ground truth annotations (y).

- $\text{FN}(y, \tilde{y}_k)$ is the number of false negatives between the predicted triplet (\tilde{y}_k) and the ground truth annotations (y).

2. **Utterance-level Reward**: Additionally, at the utterance level, a binary reward is assigned based on whether the entire set of predicted semantic triplets for the utterance matches the ground truth annotations. If all triplets in the utterance are correctly predicted, a binary reward of 1 is assigned; otherwise, a reward of 0 is assigned.

$$R_{\text{utt}} = 1 \text{ if } y = \tilde{y}_k \text{ else } 0$$

3. **Cumulative Reward**: The overall reward used for training is a combination of the triplet-level and utterance-level rewards. The cumulative reward is calculated as the sum of the triplet-level reward and the utterance-level reward:

$$R(x, y, \tilde{y}_k) = R_{\text{triplet}} + R_{\text{utt}}$$

These reward signals are then used to guide the training of the slot tagging model through policy gradient-based reinforcement learning. By maximizing the expected cumulative rewards, the system learns to improve its performance in generating accurate semantic representations of spoken utterances.

Value

1. Action-based value:

Action-based value, often denoted as $Q(s,a)$, represents the expected cumulative reward that an agent can achieve by taking a specific action a from a given state s and then following a particular policy thereafter. In this framework, action-based value estimation is crucial for determining the quality of individual actions in specific states. By approximating the action-based value function, the system can evaluate the potential outcomes of different actions and select those that are most likely to lead to favorable results. This estimation guides the agent's decision-making process during training and inference, helping it to learn optimal policies that maximize long-term rewards.

2. State-based value:

State-based value, often denoted as $V(s)$, represents the expected cumulative reward that an agent can achieve by following a particular policy from a given state s . Unlike action-based value, which focuses on the quality of individual actions, state-based value estimation provides a holistic assessment of the overall desirability of being in a particular state. State-based value estimation is essential for evaluating the potential outcomes of transitioning between different states and for guiding the agent's exploration and exploitation strategies. By approximating the state-based value function, the system can prioritize actions that lead to states associated with higher expected rewards, thus improving its decision-making efficiency.

Policy

A policy represents the strategy or decision-making process that an agent employs to select actions in different states. Policies can be broadly classified into two types: deterministic policies and stochastic policies.

1. **Deterministic Policy:** A deterministic policy directly maps states to actions without any randomness. In other words, for a given state, a deterministic policy will always output the same action. Deterministic policies are useful in situations where the environment is deterministic, and there is a clear optimal action for each state. However, in environments with uncertainty or stochasticity, deterministic policies may not be suitable as they do not allow for exploration or adaptation to changing conditions.
2. **Stochastic Policy:** A stochastic policy selects actions probabilistically, meaning that for a given state, it can output different actions with certain probabilities. Stochastic policies introduce randomness into the decision-making process, enabling the agent to explore different actions and learn more about the environment. By exploring alternative actions, stochastic policies can discover new strategies and adapt to uncertain or dynamic environments. Stochastic policies are particularly beneficial in complex and uncertain domains where the optimal action may not be immediately apparent.

In the proposed RL-based training framework for robust spoken language understanding, the choice between deterministic and stochastic policies depends on the nature of the task and the characteristics of the environment. Given the inherent uncertainty and variability in spoken language understanding tasks, a stochastic policy may be more appropriate. By allowing the agent to explore different actions probabilistically, a stochastic policy can help improve the robustness of the learned policies and enhance the agent's ability to handle diverse input scenarios, including noisy or ambiguous utterances.

Therefore, it is likely that a stochastic policy is used in the RL-based training framework described in the paper for robust spoken language understanding. This stochastic policy enables the agent to explore various strategies and adapt to different input conditions, ultimately leading to more effective and robust spoken language understanding capabilities.

Model

The approach described in the paper for robust spoken language understanding with RL-based value error recovery is primarily model-based.

1. **Model-Based Approach:** In a model-based approach, the agent builds an internal model or representation of the environment's dynamics. This model allows the agent to simulate possible future states and outcomes based on its current state and chosen actions. By leveraging this model, the agent can plan ahead and make decisions that optimize long-term rewards. Model-based approaches often require less exploration and can be more sample-efficient since they can use the learned model to generate synthetic data for planning.
2. **Model-Free Approach:** In contrast, model-free approaches directly learn a policy or value function without explicitly modeling the environment's dynamics. Instead of building an internal model, model-free approaches rely on interacting with the environment and observing the outcomes of their actions. These approaches typically involve estimating the value of states or state-action pairs directly from experience, without explicitly modeling the transitions between states. Model-free approaches are often more flexible and can handle complex environments with unknown dynamics.

In the context of the described framework, the use of reinforcement learning with a focus on training a slot tagging model and value error recovery module suggests a model-based approach. The framework incorporates an RL-based training algorithm that guides the training of the slot tagging model on ASR hypotheses with the value error recovery module. This approach involves leveraging the learned model to refine wrong values and mitigate input mismatch problems during training and testing.

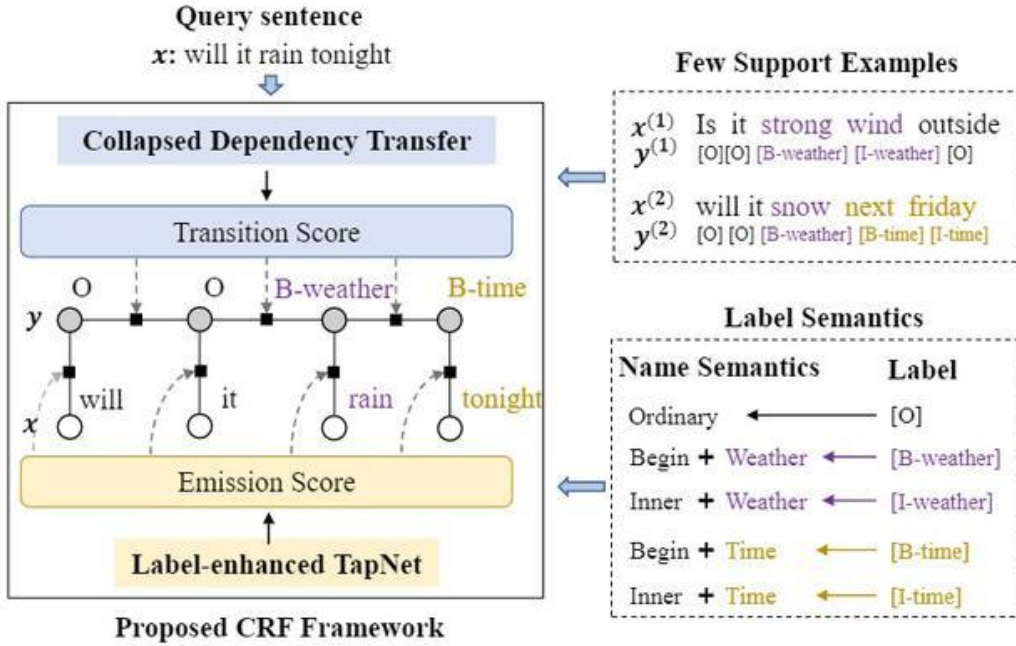


Figure 1: Our few-shot CRF framework for slot tagging.

Exploration and Exploitation

1. **Exploration:** Exploration involves trying out different actions to gain information about the environment and discover potentially better strategies. By exploring, the agent can learn more about the environment's dynamics and discover actions that may lead to higher rewards in the long run. Exploration is crucial, especially in the early stages of learning when the agent's knowledge about the environment is limited.
2. **Exploitation:** Exploitation, on the other hand, involves choosing actions that the agent believes will yield the highest immediate reward based on its current knowledge. Exploitation exploits the agent's current understanding of the environment to maximize short-term rewards. While exploitation can lead to optimal decisions based on known information, it may also prevent the agent from discovering better strategies if it relies too heavily on existing knowledge.

Balancing exploration and exploitation is a key challenge in reinforcement learning, as the agent must trade off between trying new actions to learn more about the environment (exploration) and exploiting its current knowledge to maximize immediate rewards (exploitation).

In the described framework, exploration and exploitation are likely managed through the RL-based training algorithm. During training, the agent may use exploration strategies such as ϵ -greedy or softmax exploration to try out different actions and learn about their outcomes. As the agent learns and accumulates knowledge about the environment, it gradually shifts towards exploitation, focusing more on actions that have yielded higher rewards in the past. The goal is to find a balance between exploration and exploitation to effectively learn optimal policies and maximize cumulative rewards over time.

References:

[\[2009.03095v1\] Robust Spoken Language Understanding with RL-based Value Error Recovery \(arxiv.org\)](#)

[Few-shot Slot Tagging with Collapsed Dependency Transfer and Label-enhanced \[笔记\] - 知乎 \(zhihu.com\)](#)

[The architecture of ASR hypotheses simulation. | Download Scientific Diagram \(researchgate.net\)](#)